

Development of a drone-based evaluation tool for motion analysis in athletics long jump

Studienarbeit

im Studiengang

Informationstechnik

an der Dualen Hochschule Baden-Württemberg Mannheim

von

Name, Vorname: Faust, Jakob

Abgabedatum: 16.04.2024

Bearbeitungszeitraum: 17.10.2023 - 16.04.2024

Matrikelnummer, Kurs: 5507125, TINF21IT1

Betreuer: Jürgen Schultheis

Unterschrift Betreuer: _____
Mannheim, den

Contents

List of Figures	III
Listings	IV
List of acronyms	V
1 Introduction	1
2 Tasks	3
3 Methodology	5
3.1 Software fundamentals	5
3.1.1 Programming Language and why Python	5
3.1.2 Mediapipe for detecting body poses	5
3.1.3 Why Mediapipe?	7
3.1.4 OpenCV	8
3.1.5 HDF5 file format	8
3.2 GUI development	9
3.2.1 Development framework Qt	9
4 Implementation	11
4.1 Hardware setup	11
4.1.1 Hardware selection	11
4.1.2 Hardware assembly	15
Bibliography	VII

List of Figures

3.1	Set of detectable body key points	6
3.2	HDF5 file structure	9
4.1	Frame kit	12
4.2	Motor and ESC	13
4.3	General Wiring	15
4.4	UART wiring	16
4.5	UART data frame	17

Listings

List of acronyms

AI Artificial Intelligence

BEC Battery Elimination Circuit

CPU Central Processing Unit

ESC Electronic Speed Control

FPV First Person View

GPS Global Positioning System

GPU Graphical Processing Unit

GUI Graphical User Interface

HAT Hardware Attached on Top

PDB Power Delivery board

PM Power Module

PWM Pulse Width Modulation

RPM Revolutions Per Minute

1 Introduction

Long jump is an athletic discipline that is renowned for its technical complexity and the precise movement patterns it demands from athletes. Even apparently small technical inaccuracies can significantly impact an athlete's performance. Moreover, as shown in [1] the forces during the take-off phase can reach up to 10 times the athlete's body weight, increasing the risk of serious injuries due to technical inaccuracies. Therefore, it is crucial to understand and continuously improve these movement patterns in training. However, taken the high approach velocity¹ into account, this can quickly become a difficult task. Especially the take-off phase can be very short and therefore hard to analyze.

Professional athletes often employ expensive high speed camera systems in combination with body pose markers to capture and analyze every single step they make.

Yet, this approach comes with some limitations. Due to their stationary installation, such camera systems are restricted to a fixed location. Moreover, they often combine multiple cameras like Murray et al. [2] used for sprint analysis in order to be able to capture the whole movement from the beginning of the approach until the landing. This leads to complex post-processing software requirements. Additionally, fixed markers need to be attached to an athletes body to be able to track their body position.

While these methods provide exact and reliable results, they are usually not accessible for hobby- and semi-professional athletes.

In recent years however the advances in Artificial Intelligence (AI) and especially within the area of deep neural network paved the way for analyzing methods that require less complex setups. As of 2023 deep neural networks trained for body pose detection are even used in medical applications like gait analysis [3]. Because of the already extremely high and continuously improving accuracy, its application within the area of motion analysis in long jump is treated in the scope of this work.

A semi-autonomous drone based evaluation tool is newly developed. It is supposed to offer a portable alternative to address the lack of existing opportunities in analyzing long jump performances in training. For this purpose, the drone should autonomously fly next

¹around 10 m/s in male semi-professional long jump

to the athlete throughout the whole jump, capturing their motion and therefore allow for a complete jump analysis. The drone itself is based on First Person View (FPV) drone hardware. It is build from scratch using an onboard single-board computer as flight control unit responsible for capturing the video. Additionally, a ground station software is developed to allow for a convenient jump analysis regarding the overall body pose as well as a fixed set of important parameters, i.e. knee angles, arm angles, hip position. The project's source code is available under <https://github.com/JF631/FLYJUMP>.

2 Tasks

This chapter provides an overview over all tasks that are approached in this project. Hardware related tasks are listed as well as software related tasks. Each main task is divided into several sub-tasks.

Task 1 Develop a platform independent computer software to analyze pre-recorded long-jump videos.

1. Integrate a body detection system.
2. Define an analysis output standard.
3. Implement a GUI to visualize analysis results.

Task 2 Build a quadcopter as hardware companion for recording long-jump videos. This specifically requires a camera as well as several sensors for navigation.

1. Select hardware components according defined criteria.
2. Build a flying quadcopter that is controllable from a ground. station
3. extend the quadcopter with a companion computer and a compatible camera to live stream videos.
4. test the camera accuracy.

Task 3 Implement a software that allows to control the drone from a ground station.

1. Set up a wireless connection to the quadcopter.
2. Define a communication protocol.
3. Integrate the control mechanism in the GUI.
4. Integrate a start measurement like mechanism.

Task 4 Develop a software to handle the incoming video livestream.

1. Establish a connection to the on-board companion computer.
2. Display the video frames in the GUI.
3. Offer the option for a rudimentary live analysis.

Task 5 Develop a software that can control the drone autonomously.

1. Establish between an on-board companion computer and PixHawk flight controller.
2. Implement an on-board body detection system.
3. Develop an auto-follow like controlling feature.

3 Methodology

The following chapter provides an overview over the relevant development components that are used within this project. Therefore, the used software packages are introduced before a short outline of the utilized drone hardware is given.

3.1 Software fundamentals

As the main part of this project's software will run on a portable remote computer allowing for not only to control the drone but also for performing the long jump analysis on video inputs, every software component is chosen to demand as little hardware requirements as possible. Especially no Graphical Processing Unit (GPU) is required to run the software. All image processing is performed using the Central Processing Unit (CPU) only. Furthermore, the software is designed to run platform independent.

3.1.1 Programming Language and why Python

Because of its interpreted nature and many cross-platform libraries, Python is one of the most used programming languages in the scientific area. Furthermore, it offers a high level of abstraction allowing for rapid prototyping approaches which is a key factor for this project. Besides fast implementation, Python nevertheless supports complex programming concepts like object-orientation.

Additionally, as stated in subsection 3.1.2 many machine learning and AI projects for detecting body poses have already been successfully implemented using Python.

Third party libraries and frameworks like *OpenCV* for image processing and *PyQt* for Graphical User Interface (GUI) development are widely used and therefore well documented. This leads to the decision to use Python as programming language within this work.

3.1.2 Mediapipe for detecting body poses

One of the software's main tasks is to perform a human body pose detection on video inputs. Because this part runs on the remote computer only, it can also handle pre-recorded

videos that should be evaluated.

The evaluation itself is performed using the Mediapipe framework. This framework uses a pre-trained convolutional neural network that is able to detect 33 key points in human body poses [4]. The network could theoretically even be fine-tuned to improve its accuracy on specific input types. Even if this so called *transfer-training* method requires significantly less training data compared to training a neural network from scratch, it is not applied within this project as first test runs already showed reliable results.

Furthermore, the Mediapipe framework does not require any hardware acceleration and is renowned for its precise output. Hii et al. for example showed in [3], [5] that the framework can be applied in medical gait analysis applications to replace marker based approaches. Moreover, Mediapipe offers three different detection models that differ in terms of speed and accuracy. The fastest detection model offers the lowest accuracy and vice versa.

Additionally, Mediapipe is optimized for multiple input types including videos and live streams, which is ideal for this project.

Figure 3.1 shows an overview of the 33 detectable key points.




Figure 3.1: Fixed set of detectable body key points offered by the mediapipe framework [6]

Within this work, the key points in the head area (range [0 - 10]) are not of great interest apart from visualization purposes.

The knee, hip and arm key points however will be used for angle calculations and ground contact detection. Thus, a good performance in detecting the according key points within these areas is crucial for the software's overall reliability.

3.1.3 Why Mediapipe?

In recent years many approaches towards accurate body pose detection were developed and implemented. Many of those offer decent accuracies, but often lack reasonable performance, especially when no GPU is available for hardware acceleration. Following, two common alternatives to Mediapipe, namely OpenPose and AlphaPose, are shortly presented and differentiated from the chosen Mediapipe framework.

One of the most widely used human pose detection libraries is the open source library *OpenPose*. As shown in [7] it offers a Multi-Person pose estimation that is especially useful when dealing with groups of people. However, as this project is meant to be used for long jump evaluation, only one person needs to be tracked at a time. Even though OpenPose of course can handle one person pose estimation, mediapipe outperforms OpenPose in this area. Back in 2016 Kocabas et al. achieved around 23 FPS using GPU accelerated OpenPose pose detection [8] and Osokin later proposed an improved neural network design, allowing for up to 28 FPS without hardware acceleration [9]. As of 2023 these benchmarks are still reasonable. Mediapipe however achieves speeds of up to 63% higher.

While OpenPose uses a *Bottom-Up* approach due its multi-person application, Mediapipe uses the less computational complex *Top-Down* approach.

Bottom-Up implementations first detect all body key points present in an input image and then move on to grouping the recognized points in clusters. Points in the same cluster are then assigned to one person.

Top-Down approaches however first roughly detect the overall body position within the input image and then define a region of interest¹ around the subject. The following processing therefore only needs to take this defined region of interest into account, leading to significantly less computational complexity.

AlphaPose is another open source library often used for body pose estimation. Just like OpenPose it uses a Bottom-Up approach to reliably offer multi-person body pose detection. Additionally, AlphaPose, like Mediapipe, offers multiple detection models that differ regarding accuracy and speed. Again however, Mediapipe outperforms AlphaPose because of its Top-Down approach and because AlphaPose is designed to work with GPU acceleration, rather than running on CPU only.

Another advantage of Mediapipe is its output. While OpenPose and AlphaPose offer 2D coordinates for each detected key point, Mediapipe additionally offers a depth estimation resulting in a spatial 3D coordinate for each detected key point. Thereby, more comprehensive analysis can be performed.

¹sometimes also referred to as *Bounding Box*

3.1.4 OpenCV

OpenCV is an open source library commonly used for image processing in the area of computer vision and machine learning. It is written in C++, thus offering high performance in numerical operations, especially matrix operations.

`python-opencv` is the python wrapper for OpenCV which will be used in this project to efficiently read and process video frames. The python wrapper imports the underlying C++ functions as modules to take advantage of C++'s efficiency, resulting in significantly higher performance compared to equivalent Python only implementations. Furthermore, it is fully compatible with the `numpy` library, which allows for seamless conversion between numpy arrays and OpenCV image matrices.

3.1.5 HDF5 file format

To avoid analyzing the same jump, therefore the same video file, multiple times, the jump parameters should be saved after the analysis process alongside with the annotated video file, that visualizes the detected body pose (see Figure 3.1). Because parameters such as knee angles, arm angles, ground contact, etc. need to be stored frame-wise, a structured file format is suitable.

One common open source structured file format is HDF5. It is an acronym for Hierarchical Data Format (Version 5) and is very helpful to store large amounts of data as well as heterogeneous data. As the name already suggests, data is stored in a hierarchical, tree-like way.

A HDF5 tree mainly consists of three pre-defined components that allow to organize data in a file-system like fashion. The tree root, groups and datasets. *Groups* are folder-like structures that can either contain more groups or Datasets. *Datasets* hold the actual data that need to be stored.

Furthermore, each level (tree root, groups and datasets) can hold additional information via metadata. The metadata could for example contain information about metrics, timestamps or any other describing information. Therefore, each HDF5 file itself becomes a self-describing file that does not require any more than the included information to be interpreted correctly.




Figure 3.2: Principle HDF5 file structure ²

Within this work, one HDF5 file will be created per jump analysis. The frames are stored in a group each and the actual calculated parameters are saved group-wise as datasets.

3.2 GUI development

The software that is developed within this work should be useable on-field to allow for a fast analysis process. Besides the video analysis, the drone control should be embedded in the same software to always guarantee control over the drone.

Thus, a simple GUI is developed to offer a convenient drone control and video analysis process.

The GUI is developed using Qt and its Python binding PyQt. Both are presented in this section.

3.2.1 Development framework Qt

Qt is a development framework based on the programming language C++. It includes a GUI toolkit and therefore enables platform-independent application development. All common platforms including Linux, Windows and MacOS are supported by Qt. Additionally, both mobile operating systems, Android and iOS, are supported.

This project however will focus on the development of an application that is able to run under Windows, Linux and MacOS.

The Qt framework is mainly chosen because of its rich and comprehensive documentation³ and the availability of the well-supported Python binding *PyQt*.

As Qt is based on C++ all Qt source files are translated to C++ code. This step is

³<https://doc.qt.io/>

realized by the **Meta Object Compiler (MOC)** which is integrated as pre-processor. Thus, all Qt files are translated to so-called *Meta Object Code*, which can be seen as C++ source code with some enhancements. The most important enhancement is the signal and slot functionality which allows for an easy communication between different software and design elements (e.g. show a message dialog when a button is clicked).

Another important enhancement for this work is the convenient multi-threading management necessary for offering a responsive GUI even when cpu-bound calculations such as image processing is performed.

The GUI module PyQt

The discrepancy between Qt as C++ based framework and Python as chosen programming language for this project (as explained in ??) can be overcome using Qt's Python binding PyQt. By using PyQt we can combine Python's machine learning advantages with Qt's platform-independent GUI development. More specifically *PyQt5* ⁴ is used.

It allows building complex Qt applications using Python only instead of C++. All other described advantages that Qt offer remain valid despite the use of PyQt. Thus, the whole software within this project including the GUI can be developed using Python as programming language.

⁴<https://pypi.org/project/PyQt5/>

4 Implementation

This chapter focuses on the overall project's implementation. It mainly covers the four parts hardware assembly, long jump analysis, drone control and their consolidation into one GUI.

4.1 Hardware setup

In order to capture high-quality video recordings that cover a complete long jump, from the first step all the way to the landing, a drone is used to fly next to the athlete throughout the whole process. Thus, a drone in form of a quadcopter is built from scratch. Its control will be integrated seamlessly in the projects' GUI.

This section introduces the hardware components that are used for building this drone as well as its flight control unit.

A short outline of the hardware is given in subsection 4.1.1, while subsection 4.1.2 focuses on the overall assembly of the selected hardware.

4.1.1 Hardware selection

Currently, commercial drone hardware on the market is mainly separable into the two large areas of fully remote controlled FPV hardware and hardware for (autonomous) drones that can usually carry more load, e.g. heavy cameras. Even though the quadcopter in this project needs to be remotely controllable from a ground station pc, it is still more likely to be located in the latter one.

Generally the hardware was chosen based on the following criteria:

- price
- compatibility
- size

Flight Hardware

The main hardware that a quadcopter needs to fly will, in the following, be referred to as *flight hardware*. This includes frame, motors, rotors, Electronic Speed Controls (ESCs) and a Power Delivery board (PDB).

The main platform on which all drone hardware is mounted, is referred to as a quadcopter's frame. As this project's drone does not need to carry any heavy load, such as high precision camera systems or other sensors, a rather compact frame would theoretically be sufficient. However, compact frames tend to be less stable compared to larger frame sizes which could lead to a lower video recording quality and thus require more complex post-processing software. Moreover, the assembly process on larger frames is more convenient and replacing parts is easier. Additionally, compact frames are most commonly used in areas that demand quick reaction times for high speed flight maneuvers, e.g. in drone racing. This however is not needed in this project's context.

Taken the mentioned considerations into account the mid-sizes *Holybro S500 V2* frame kit is chosen. Besides the frame, the kit also includes a landing gear and rotors. Moreover, the main platform includes a PDB to split the battery's power equally to all four motors. An overview of all included parts is given in Figure 4.1.





Figure 4.1: Holybro S500 V2 frame kit

Besides the frame, motors and compatible ESCs are crucial flight hardware components. Each motor requires an own ESC that translates signals from a flight control unit to a voltage and thereby control the motors' rotation speed. To guarantee compatibility, both components were chosen from Holybro as well and can be seen in Figure 4.2.



(a) 920KV Motor



(b) Electronic Speed Control

Figure 4.2: Motor (a) and ESC (b)

The drones' motors performance capabilities are defined by the number of Revolutions Per Minute (RPM) they can perform per 1V input. As can be seen in Figure 4.2a, this link between rotation speed and input voltage is expressed in the arbitrary unit *KV*. The chosen motors are capable of rotating with a speed of 920 RPM per 1V input voltage. Put into context, this is a common rotation speed in commercial and hobby drone applications. Racing drones however, operate at motor speeds of up to 3500 KV.

Control Hardware

In order to perform flight maneuvers with a quadcopter, each motor must be controllable individually. The calculation of the correct rotation speeds is generally performed by a *flight control unit*. Usually, it receives directional instructions from a remote control as input, combines them with many parameters (e.g. GPS position, height over ground, speed, etc.) and generates a (PWM) output signal for each motor.

Within this project the flight controller needs to deal with two different inputs. First, the ground station which can be seen as a remote control in this case. Additionally, the drone should be able to fly autonomously next to an athlete during their long jump training. Here, the second input gets important. The autonomous fly option requires the quadcopter to perform a person detection and therefore image processing on-board. As the flight controller itself is not able to perform such calculations, an additional *companion computer* is required. This companion computer will then send directional instructions just like the ones from the ground station to the flight controller and thereby control the drone.

The combination of flight controller and on-board companion computer will in the following be called *control hardware*.

There are many types of different flight controllers available commercially. However, most of them are not meant to be used in combination with a companion computer.

Two of the most commonly used flight controllers in autonomous drone projects are the *PixHawk* and the *Navio2*. They are often chosen because they both work together seamlessly with a companion computer. The former is a totally independent system which can also operate without any supporting computer. The latter is implemented as Hardware Attached on Top (HAT) specifically designed for a Raspberry Pi. Thus, it does not include an own CPU but uses the Raspberry Pi's resources to perform flight relevant calculations. A detailed comparison between both flight controllers is given in Table 4.1.

Flight Controller Comparison		
Criteria	PixHawk	Navio2
Processor	ARM Cortex M4 with FPU / 32-bit co-processor	Depends on Raspberry Pi version
Sensors	ST Micro 16-bti gyroscope, ST Micro 14-bit accelerometer, MEAS barometer	MPU9250 9DOF IMU, LSM9DS1 9DOF IMU, MS5611 Barometer, U-blox M8N Glonass/GPS/Beidou
Interfaces	UART, Spektrum DSM, PPM / S.BUS input, I2C, SPI, CAN, USB, 3.3V and 6.6V ADC input, 8 PWM outputs, 6 Auxiliary outputs	UART, I2C, ADC, PPM / S.BUS input, 14 PWM outputs
Dimensions (W x H x L) in mm	50 × 15.5 × 81.5	55 × 65
Other	Failsafe options (e.g. extra power supply, GPS, etc.)	None
Price (Eur.)	TBD	TBD

Table 4.1: Comparison between PixHawk and Navio2 flight controller.

As can be seen, both flight controllers offer different interfaces to connect additional hardware. Moreover, both systems include sensors, mostly to gather information about drone's current position and inertia. Here, the Navio2 even offers more sensors, as it already includes a Global Positioning System (GPS) sensor, while the PixHawk relies on an external one.

For this project, the PixHawk was chosen over the Navio2 mainly for three reasons.

First, it is on the market for a long time already and thus have a large community support. Secondly, as it is an independent system, a failure of the companion computer will not lead to a crash. Lastly, it allows for a wide range of companion computers, while the Navio2 can only interoperate with a Raspberry Pi.

Furthermore, the mentioned considerations lead to easy rapid prototyping approaches, as the drone can be manually flown without a companion computer in a first implementation.

4.1.2 Hardware assembly

In the following, a high-level overview of the quadcopters' hardware setup is given. The general wiring is shown and explained before a short introduction of some important communication protocols is given.

General wiring

In the following Figure 4.3 the general wiring layout is shown. The whole system is powered from one power source only.




Figure 4.3: General wiring of the hardware setup.

Power connections are labeled red, Blue connections are used for Pulse Width Modulation (PWM) signals, Green connections are serial connections, orange connections are more specifically serial UART connections.

This results in some challenges in providing the correct voltage for each connected device. In the current setup this task is taken over by three devices. The Power Module (PM)

is directly connected to the battery. It transfers the battery's voltage to the PDB and a lower 5V voltage to the PixHawk flight controller. The PDB itself is a parallel circuit, thus providing the same voltage (battery voltage) to each output. The third device is a Battery Elimination Circuit (BEC) which is directly connected to the PDB and delivers a constant 5V output. This can be used to power a companion computer such as a RaspberryPi.

All other required peripherals are powered by the PixHawk flight controller itself. The main peripherals used in this project are a telemetry module which is used for communication with a ground station and a GPS module used for improving the drones capabilities to follow a defined trajectory, which is specifically useful for auto-return and landing features. Two more peripherals, a buzzer to output audio warning signals and a manual kill switch which can immediately stop all four motors, are installed mainly for safety reasons.

The hardware components that actually control the motor rotation speeds, the ESCs, are connected to the PDB for power supply as well as to the flight controller that calculates the correct rotation speeds based on the wanted flight maneuvers and outputs a PWM signal for each motor.

The presented overall wiring is rather complex but allows relying on one power source only instead of using multiple power sources for flight hardware and control hardware including peripherals respectively.

Communication between components

The drone setup needs hardware components to communicate with each other in order to transfer control signals from either the companion computer or from the ground station to the flight controller. Even if both options origin from different sources, they use the same device-to-device communication protocol. The protocol used for this purpose is the *UART* protocol, which is acronym for universal asynchronous receiver / transmitter protocol. It is based on a serial, full duplex connection using six connections. The connection layout is shown in Figure 4.4,




Figure 4.4: UART wiring

where RTS / CTS denote Ready to send and Clear to send. RX / TX represent the actual data receiving and sending connections respectively.

UART transfers data using data frames with minimal overhead. A typical UART frame consists of just a start bit, data bits, a parity bit and a stop bit. Figure 4.5 visualizes such a data frame.




Figure 4.5: Exmple of an UART data frame. The blue marked area is the actual data part that is transmitted.

The shown UART data frame includes 5 data bits, however, the amount can vary between 5 and 9. Moreover, the included even parity bit is optional. The idle state is set to a voltage that represents HIGH level on purpose, so that any connection failure is easily detectable. UART can work with any voltages to denote HIGH and LOW levels. In the quadcopter setup HIGH is represented by 5V and LOW by GND.

Bibliography

- [1] A. Seyfarth et al. “Dynamics of the long jump”. In: *Journal of Biomechanics* 32.12 (1999), pp. 1259–1267. ISSN: 0021-9290. DOI: [https://doi.org/10.1016/S0021-9290\(99\)00137-2](https://doi.org/10.1016/S0021-9290(99)00137-2). URL: <https://www.sciencedirect.com/science/article/pii/S0021929099001372>.
- [2] Murray Evans et al. “Automatic high fidelity foot contact location and timing for elite sprinting”. In: *Machine Vision and Applications* 32.05 (2021). DOI: 10.1007/s00138-021-01236-z.
- [3] Chang Soon Tony Hii et al. “Marker Free Gait Analysis using Pose Estimation Model”. In: *2022 IEEE 20th Student Conference on Research and Development (SCoReD)*. 2022, pp. 109–113. DOI: 10.1109/SCoReD57082.2022.9974096.
- [4] Valentin Bazarevsky et al. “BlazePose: On-device Real-time Body Pose tracking”. In: *CoRR* abs/2006.10204 (2020). arXiv: 2006.10204. URL: <https://arxiv.org/abs/2006.10204>.
- [5] Amit Gupta et al. “Knee Flexion/Extension Angle Measurement for Gait Analysis Using Machine Learning Solution “MediaPipe Pose” and Its Comparison with Kinovea®”. In: *IOP Conference Series: Materials Science and Engineering* 1279.1 (March 2023), p. 012004. DOI: 10.1088/1757-899X/1279/1/012004. URL: <https://dx.doi.org/10.1088/1757-899X/1279/1/012004>.
- [6] Google Developers. *Pose landmark detection guide*. 2023. URL: https://developers.google.com/mediapipe/solutions/vision/pose_landmarker (visited on November 5, 2023).
- [7] Zhe Cao et al. “Realtime Multi-Person 2D Pose Estimation using Part Affinity Fields”. In: *CoRR* abs/1611.08050 (2016). arXiv: 1611.08050. URL: <http://arxiv.org/abs/1611.08050>.
- [8] Muhammed Kocabas et al. “MultiPoseNet: Fast Multi-Person Pose Estimation using Pose Residual Network”. In: *Proceedings of the European Conference on Computer Vision (ECCV)*. September 2018.

-
- [9] Daniil Osokin. “Real-time 2D Multi-Person Pose Estimation on CPU: Lightweight OpenPose”. In: *CoRR* abs/1811.12004 (2018). arXiv: 1811.12004. URL: <http://arxiv.org/abs/1811.12004>.